

Department of Electronics & Telecommunication
Engineering
University of Moratuwa

EN3150 - Pattern Recognition



Learning from data and related challenges and linear
models for regression

EN3150 Assignment 01

Name: Dilshan N.L.
Index No: 210129P

Date - 2024.09.02

Contents

1 Data Pre-Processing	2
2 Learning from Data	3
1. Generating Data Using Listing 1	3
2. Running Listing 2 and Observing Training and Testing Data	3
3. Fitting Linear Regression Model and Observing Different Instances	4
4. Increasing the Number of Data Samples to 10,000	5
3 Linear regression on real world data	7
1. Loading the Dataset	7
2. Independent and Dependent Variables	7
3. Is it possible to apply linear regression?	9
4. Handling NaN/Missing Values	10
5. and 6. Selecting Features and Splitting Data	11
7. Training a Linear Regression Model	12
8. Identifying the Most Contributing Variable	13
9. Additional Feature Selection and Model Training	13
10. Calculating Statistical Measures	14
11. Significant and Insignificant features	15
4 Performance evaluation of Linear regression	16
5 Linear regression impact on outliers	17
2. What happens when $a \rightarrow 0$?	17
3. Minimizing the influence of data points with $ r_i \geq 40$	17

1 Data Pre-Processing

Max-abs scaling is the preferred scaling method for the given features.

Reason: Max-Abs Scaling is ideal for this scenario because it scales the feature values relative to their maximum absolute value while preserving zero values. This method ensures that the zero values remain unchanged, which is crucial for maintaining the structure of the data if zeroes are meaningful in the feature. In contrast, Standard Scaling and Min-Max Scaling would shift the zero values or alter their meaning, which could be undesirable if preserving the original structure is important.




Figure 1: Feature values of a dataset.

2 Learning from Data

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import pandas as pd
import statsmodels.api as sm
```

1. Generating Data Using Listing 1

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

# Generate 100 samples
n_samples = 100

# Generate X values (uniformly distributed between 0 and 10)
X = 10 * np.random.rand(n_samples, 1)

# Generate epsilon values (normally distributed with mean 0 and standard deviation 15)
epsilon = np.random.normal(0, 15, n_samples)

# Generate Y values using the model  $Y = 3 + 2 * X + \epsilon$ 
Y = 3 + 2 * X + epsilon[:, np.newaxis]
```


2. Running Listing 2 and Observing Training and Testing Data

```
r = np.random.randint(104)
# Split the data into training and test sets (80% train, 20% test)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=r)

# Plot the data points
plt.figure(figsize=(10, 6))
plt.scatter(X_train, Y_train, alpha=1, marker='o', color='red', label='Training Data')
plt.scatter(X_test, Y_test, alpha=1, marker='s', color='blue', label='Testing Data')
plt.show()
```

Observation:


- Each time we run the code, the training and testing data will be different. This is because the 'random_state' used in 'train_test_split' is generated using a random integer('r'), which changes on every run.
- **Reason:** The 'random_state' controls the shuffling of data before splitting into training and test sets. Since 'r' changes each time, the split is different in each run.



3. Fitting Linear Regression Model and Observing Different Instances

```
for i in range(10):
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=np.random.)
    model = LinearRegression()
    model.fit(X_train, Y_train)
    Y_pred_train = model.predict(X_train)
    plt.plot(X_train, Y_pred_train, label=f'LR {i+1}')

plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.show()
```



Observation:

- The linear regression model will vary slightly between each instance.
- Reason: Each time the data is split differently due to the changing 'random_state', the training data the model learns from is different. This leads to slight variations in the fitted model.

4. Increasing the Number of Data Samples to 10,000

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

# Generate 100 samples
n_samples = 10000

# Generate X values (uniformly distributed between 0 and 10)
X = 10 * np.random.rand(n_samples, 1)

# Generate epsilon values (normally distributed with mean 0 and standard deviation 15)
epsilon = np.random.normal(0, 15, n_samples)

# Generate Y values using the model  $Y = 3 + 2 * X + \text{epsilon}$ 
Y = 3 + 2 * X + epsilon[:, np.newaxis]



r = np.random.randint(104)
# Split the data into training and test sets (80% train, 20% test)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=r)

# Plot the data points
plt.figure(figsize=(10, 6))
plt.scatter(X_train, Y_train, alpha=1, marker='o', color='red', label='Training Data')
plt.scatter(X_test, Y_test, alpha=1, marker='s', color='blue', label='Testing Data')
plt.show()

for i in range(10):
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=np.random.
    model = LinearRegression()
    model.fit(X_train, Y_train)
    Y_pred_train = model.predict(X_train)
    plt.plot(X_train, Y_pred_train, label=f'LR {i+1}')

plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.show()

```

**Observation:**

- When the number of samples is increased to 10000, the linear regression model instances will exhibit much less variation compared to when there are only 100 samples.
- Reason:
 - Larger Dataset: With a larger dataset, the training data becomes more representative of the entire data distribution. Thus, even with different random splits, the model tends to converge towards a more consistent fit.
 - Reduced Impact of Random Variations: Random variations in the training data have less influence when there are more samples, leading to more stable models across different instances.

3 Linear regression on real world data

1. Loading the Dataset

```
# If package not installed, install it using pip install ucimlrepo
from ucimlrepo import fetch_ucirepo

# fetch dataset
infrared_thermography_temperature = fetch_ucirepo(id=925)

# data (as pandas dataframes)
X = infrared_thermography_temperature.data.features
y = infrared_thermography_temperature.data.targets

# metadata
print(infrared_thermography_temperature.metadata)

# variable information
print(infrared_thermography_temperature.variables)
```

2. Independent and Dependent Variables

- Independent variables: These are the features in X.
- Dependent variables: These are the target values in y.

```
print(infrared_thermography_temperature.data)
```

```
{'ids':      SubjectID
0      161117-1
1      161117-2
2      161117-3
3      161117-4
4      161117-5
...      ...
1015   180425-05
1016   180425-06
1017   180502-01
1018   180507-01
1019   180514-01
```

```
[1020 rows x 1 columns], 'features':      Gender      Age      Ethnicity  T_atm  Humidity
0      Male  41-50      White  24.0      28.0      0.8
1    Female  31-40  Black or African-American  24.0      26.0      0.8
2    Female  21-30      White  24.0      26.0      0.8
3    Female  21-30  Black or African-American  24.0      27.0      0.8
4      Male  18-20      White  24.0      27.0      0.8
...      ...      ...      ...      ...      ...
1015  Female  21-25      Asian  25.7      50.8      0.6
1016  Female  21-25      White  25.7      50.8      0.6
1017  Female  18-20  Black or African-American  28.0      24.3      0.6
1018   Male  26-30      Hispanic/Latino  25.0      39.8      0.6
1019  Female  18-20      White  23.8      45.6      0.6
```

```
      T_offset1  Max1R13_1  Max1L13_1  aveAllR13_1  ...  T_FHCC1  T_FHRC1  \
0      0.7025    35.0300    35.3775    34.4000    ...  33.5775  33.4775
1      0.7800    34.5500    34.5200    33.9300    ...  34.0325  34.0550
2      0.8625    35.6525    35.5175    34.2775    ...  34.9000  34.8275
3      0.9300    35.2225    35.6125    34.3850    ...  34.4400  34.4225
4      0.8950    35.5450    35.6650    34.9100    ...  35.0900  35.1600
...      ...      ...      ...      ...      ...  ...      ...
```


1015	1.2225	35.6425	35.6525	34.8575	...	35.1075	35.3475
1016	1.4675	35.9825	35.7575	35.4275	...	35.3100	35.2175
1017	0.1300	36.4075	36.3400	35.8700	...	35.4350	35.2400
1018	1.2450	35.8150	35.5250	34.2950	...	34.8400	35.0200
1019	0.8675	35.7075	35.5825	34.8875	...	34.5475	34.6500

	T_FHLC1	T_FHBC1	T_FHTC1	T_FH_Max1	T_FHC_Max1	T_Max1	T_OR1	\
0	33.3725	33.4925	33.0025	34.5300	34.0075	35.6925	35.6350	
1	33.6775	33.9700	34.0025	34.6825	34.6600	35.1750	35.0925	
2	34.6475	34.8200	34.6700	35.3450	35.2225	35.9125	35.8600	
3	34.6550	34.3025	34.9175	35.6025	35.3150	35.7200	34.9650	
4	34.3975	34.6700	33.8275	35.4175	35.3725	35.8950	35.5875	
...	
1015	35.4000	35.1375	35.2750	35.8525	35.7475	36.0675	35.6775	
1016	35.2200	35.2075	35.0700	35.7650	35.5525	36.5000	36.4525	
1017	35.2275	35.3675	35.3425	36.3750	35.7100	36.5350	35.9650	
1018	34.9250	34.7150	34.5950	35.4150	35.3100	35.8600	35.4150	
1019	34.6700	34.2150	34.7100	35.1525	35.1175	35.9725	35.8900	

	T_OR_Max1
0	35.6525
1	35.1075
2	35.8850
3	34.9825
4	35.6175
...	...
1015	35.7100
1016	36.4900
1017	35.9975
1018	35.4350
1019	35.9175

[1020 rows x 33 columns], 'targets':			aveOralF	aveOralM
0	36.85	36.59		
1	37.00	37.19		
2	37.20	37.34		
3	36.85	37.09		
4	36.80	37.04		
...		
1015	36.95	36.99		
1016	37.25	37.19		
1017	37.35	37.59		
1018	37.15	37.29		
1019	37.05	37.19		

[1020 rows x 2 columns], 'original':				SubjectID	aveOralF	aveOralM	Gender	Age
0	161117-1	36.85	36.59	Male	41-50		White	
1	161117-2	37.00	37.19	Female	31-40	Black or African-American		
2	161117-3	37.20	37.34	Female	21-30		White	
3	161117-4	36.85	37.09	Female	21-30	Black or African-American		
4	161117-5	36.80	37.04	Male	18-20		White	
...	
1015	180425-05	36.95	36.99	Female	21-25		Asian	
1016	180425-06	37.25	37.19	Female	21-25		White	
1017	180502-01	37.35	37.59	Female	18-20	Black or African-American		
1018	180507-01	37.15	37.29	Male	26-30		Hispanic/Latino	
1019	180514-01	37.05	37.19	Female	18-20		White	

	T_atm	Humidity	Distance	T_offset1	...	T_FHCC1	T_FHRC1	T_FHLC1	\
0	24.0	28.0	0.8	0.7025	...	33.5775	33.4775	33.3725	
1	24.0	26.0	0.8	0.7800	...	34.0325	34.0550	33.6775	
2	24.0	26.0	0.8	0.8625	...	34.9000	34.8275	34.6475	
3	24.0	27.0	0.8	0.9300	...	34.4400	34.4225	34.6550	
4	24.0	27.0	0.8	0.8950	...	35.0900	35.1600	34.3975	
...	
1015	25.7	50.8	0.6	1.2225	...	35.1075	35.3475	35.4000	
1016	25.7	50.8	0.6	1.4675	...	35.3100	35.2175	35.2200	
1017	28.0	24.3	0.6	0.1300	...	35.4350	35.2400	35.2275	
1018	25.0	39.8	0.6	1.2450	...	34.8400	35.0200	34.9250	
1019	23.8	45.6	0.6	0.8675	...	34.5475	34.6500	34.6700	

	T_FHBC1	T_FHTC1	T_FH_Max1	T_FHC_Max1	T_Max1	T_OR1	T_OR_Max1
0	33.4925	33.0025	34.5300	34.0075	35.6925	35.6350	35.6525
1	33.9700	34.0025	34.6825	34.6600	35.1750	35.0925	35.1075
2	34.8200	34.6700	35.3450	35.2225	35.9125	35.8600	35.8850
3	34.3025	34.9175	35.6025	35.3150	35.7200	34.9650	34.9825
4	34.6700	33.8275	35.4175	35.3725	35.8950	35.5875	35.6175
...
1015	35.1375	35.2750	35.8525	35.7475	36.0675	35.6775	35.7100
1016	35.2075	35.0700	35.7650	35.5525	36.5000	36.4525	36.4900
1017	35.3675	35.3425	36.3750	35.7100	36.5350	35.9650	35.9975
1018	34.7150	34.5950	35.4150	35.3100	35.8600	35.4150	35.4350
1019	34.2150	34.7100	35.1525	35.1175	35.9725	35.8900	35.9175

```
[1020 rows x 36 columns], 'headers': Index(['SubjectID', 'aveOralF', 'aveOralM', 'Gender', 'Age', 'E',
      'T_atm', 'Humidity', 'Distance', 'T_offset1', 'Max1R13_1', 'Max1L13_1',
      'aveAllR13_1', 'aveAllL13_1', 'T_RC1', 'T_RC_Dry1', 'T_RC_Wet1',
      'T_RC_Max1', 'T_LC1', 'T_LC_Dry1', 'T_LC_Wet1', 'T_LC_Max1', 'RCC1',
      'LCC1', 'canthiMax1', 'canthi4Max1', 'T_FHCC1', 'T_FHRC1', 'T_FHLC1',
      'T_FHBC1', 'T_FHTC1', 'T_FH_Max1', 'T_FHC_Max1', 'T_Max1', 'T_OR1',
      'T_OR_Max1'],
      dtype='object')}]
```

```
print(f"Number of Independent Variables: {X.shape[1]}")
print(f"Number of Dependent Variables: {y.shape[1]}")
print(X,y)
```

```
Number of Independent Variables: 33
Number of Dependent Variables: 2
```

3. Is it possible to apply linear regression?

In this dataset, we have non-numeric data such as age ranges, sex, and other categorical variables. To apply linear regression to these types of data, they need to be converted into a numerical format. This can be achieved using the following methods:

1. **Label Encoding:** Assigns a unique integer to each category. This is suitable for ordinal data where categories have a meaningful order, such as 'low', 'medium', and 'high'.
2. **One-Hot Encoding:** Creates binary columns for each category, indicating the presence or absence of each category. This method is ideal for nominal data without an inherent order, such as 'sex' or 'ethnicity'.
3. **Ordinal Encoding:** Assigns integer values to categories based on their inherent order. This method is appropriate for ordinal variables where the sequence of categories carries significance, such as 'age ranges'.
4. **Binning:** Converts continuous variables into discrete categories or bins. This is useful for grouping continuous data, like 'age', into meaningful ranges.

By employing these encoding techniques, non-numeric data can be effectively transformed into a numerical format suitable for linear regression analysis.

4. Handling NaN/Missing Values

The provided code is not correct. Because we must remove both the X and y values corresponding to a missing value.

`table.dropna()` ensures that we remove rows with any missing values across the entire dataset, maintaining consistency and alignment. `X.dropna()` and `y.dropna()` separately might lead to mismatched data and additional complexity, especially when dealing with feature and target data.

```
import pandas as pd

table = pd.concat([X, y], axis = 1)
# Count missing values for each column
missing_values_per_column = table.isnull().sum()
print("Missing values per column:")
print(missing_values_per_column)
# Count the total number of missing values in the DataFrame
total_missing_values = table.isnull().sum().sum()
print(f"Total number of missing values in the DataFrame: {total_missing_values}")
```

Missing values per column:

Gender	0
Age	0
Ethnicity	0
T_atm	0
Humidity	0
Distance	2
T_offset1	0
Max1R13_1	0
Max1L13_1	0
aveAllR13_1	0
aveAllL13_1	0
T_RC1	0
T_RC_Dry1	0
T_RC_Wet1	0
T_RC_Max1	0
T_LC1	0
T_LC_Dry1	0
T_LC_Wet1	0
T_LC_Max1	0
RCC1	0
LCC1	0
canthiMax1	0
canthi4Max1	0
T_FHCC1	0
T_FHRC1	0
T_FHLC1	0
T_FHBC1	0
T_FHTC1	0
T_FH_Max1	0
T_FHC_Max1	0
T_Max1	0
T_OR1	0
T_OR_Max1	0
aveOralF	0
aveOralM	0

dtype: int64

```

Total number of missing values in the DataFrame: 2
table = table.dropna()
# Count missing values for each column
missing_values_per_column = table.isnull().sum()
print("Missing values per column:")
print(missing_values_per_column)
# Count the total number of missing values in the DataFrame
total_missing_values = table.isnull().sum().sum()
print(f"Total number of missing values in the DataFrame: {total_missing_values}")

Missing values per column:
Gender          0
Age             0
Ethnicity       0
T_atm           0
Humidity        0
Distance        0
T_offset1       0
Max1R13_1       0
Max1L13_1       0
aveAllR13_1     0
aveAllL13_1     0
T_RC1           0
T_RC_Dry1       0
T_RC_Wet1       0
T_RC_Max1       0
T_LC1           0
T_LC_Dry1       0
T_LC_Wet1       0
T_LC_Max1       0
RCC1            0
LCC1            0
canthiMax1      0
canthi4Max1     0
T_FHCC1         0
T_FHRC1         0
T_FHLC1         0
T_FHBC1         0
T_FHTC1         0
T_FH_Max1       0
T_FHC_Max1      0
T_Max1          0
T_OR1           0
T_OR_Max1       0
aveOralF        0
aveOralM        0
dtype: int64
Total number of missing values in the DataFrame: 0

```

5. and 6. Selecting Features and Splitting Data

```

# Selecting 'aveOralM' as the dependent variable
y = y[['aveOralM']]

# Selecting 'Age' and four other features based on preference
X = X[['Age', 'T_OR1', 'T_OR_Max1', 'T_FHC_Max1', 'T_FH_Max1']]

print(X,y)

```

```
# Splitting the data
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

	Age	T_OR1	T_OR_Max1	T_FHC_Max1	T_FH_Max1
0	41-50	35.6350	35.6525	34.0075	34.5300
1	31-40	35.0925	35.1075	34.6600	34.6825
2	21-30	35.8600	35.8850	35.2225	35.3450
3	21-30	34.9650	34.9825	35.3150	35.6025
4	18-20	35.5875	35.6175	35.3725	35.4175
...
1015	21-25	35.6775	35.7100	35.7475	35.8525
1016	21-25	36.4525	36.4900	35.5525	35.7650
1017	18-20	35.9650	35.9975	35.7100	36.3750
1018	26-30	35.4150	35.4350	35.3100	35.4150
1019	18-20	35.8900	35.9175	35.1175	35.1525

```
[1020 rows x 5 columns]      aveOralM
0      36.59
1      37.19
2      37.34
3      37.09
4      37.04
...      ...
1015    36.99
1016    37.19
1017    37.59
1018    37.29
1019    37.19

[1020 rows x 1 columns]
```

7. Training a Linear Regression Model

```
print(X.columns)
Index(['Age', 'T_OR1', 'T_OR_Max1', 'T_FHC_Max1', 'T_FH_Max1'], dtype='object')

print(X.Age)
0      41-50
1      31-40
2      21-30
3      21-30
4      18-20
...
1015    21-25
1016    21-25
1017    18-20
1018    26-30
1019    18-20
Name: Age, Length: 1020, dtype: object

def convert_age_range(age_range):
    """Converts the age range to a single average value"""
    if '>' in age_range:
        return int(age_range.replace('>', '').strip())
    lower, upper = map(int, age_range.split('-'))
    return (lower + upper) / 2
```

```
X.Age = X.Age.apply(convert_age_range)
print(X.Age)
```

```
0      45.5
1      35.5
2      25.5
3      25.5
4      19.0
```

```
...
1015    23.0
1016    23.0
1017    19.0
1018    28.0
1019    19.0
```

```
Name: Age, Length: 1020, dtype: float64
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_27752\2353021610.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/index.html

```
X.Age = X.Age.apply(convert_age_range)
```

```
from sklearn.linear_model import LinearRegression
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model = LinearRegression()
model.fit(X_train, y_train)
```

```
# Coefficients corresponding to independent variables
```

```
coefficients = model.coef_
print(f"Estimated Coefficients: {coefficients}")
```

```
Estimated Coefficients: [[ 0.00113644  0.05647584  0.49937613 -0.08398371  0.36994022]]
```

8. Identifying the Most Contributing Variable

The variable with the highest absolute value in the coefficient array contributes the most:

```
import numpy as np
```

```
max_contributor_index = np.argmax(np.abs(coefficients))
most_contributing_feature = X.columns[max_contributor_index]
print(f"Most contributing feature: {most_contributing_feature}")
```

```
Most contributing feature: T_OR_Max1
```

9. Additional Feature Selection and Model Training

```
X = X[['T_OR1', 'T_OR_Max1', 'T_FHC_Max1', 'T_FH_Max1']]
print(X)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model.fit(X_train, y_train)
coefficients = model.coef_
print(f"Estimated Coefficients: {coefficients}")
```

```
      T_OR1  T_OR_Max1  T_FHC_Max1  T_FH_Max1
0      35.6350      35.6525      34.0075      34.5300
1      35.0925      35.1075      34.6600      34.6825
2      35.8600      35.8850      35.2225      35.3450
```

3	34.9650	34.9825	35.3150	35.6025
4	35.5875	35.6175	35.3725	35.4175
...
1015	35.6775	35.7100	35.7475	35.8525
1016	36.4525	36.4900	35.5525	35.7650
1017	35.9650	35.9975	35.7100	36.3750
1018	35.4150	35.4350	35.3100	35.4150
1019	35.8900	35.9175	35.1175	35.1525

[1020 rows x 4 columns]

Estimated Coefficients: [[0.09199696 0.4640698 -0.08733171 0.37088645]]

10. Calculating Statistical Measures

```
from sklearn.metrics import mean_squared_error
```

```
# Residual sum of squares (RSS)
```

```
y_pred = model.predict(X_test)
```

```
RSS = np.sum(np.square(y_test - y_pred))
```

```
# Residual Standard Error (RSE)
```

```
N = len(y_test)
```

```
d = X_train.shape[1]
```

```
RSE = np.sqrt(RSS / (N - d - 1))
```

```
# Mean Squared Error (MSE)
```

```
MSE = mean_squared_error(y_test, y_pred)
```

```
# R-squared statistic
```

```
R_squared = model.score(X_test, y_test)
```

```
# Standard Error, t-statistic, p-value
```

```
import statsmodels.api as sm
```

```
X_train_with_const = sm.add_constant(X_train)
```

```
ols_model = sm.OLS(y_train, X_train_with_const).fit()
```

```
standard_errors = ols_model.bse
```

```
t_statistics = ols_model.tvalues
```

```
p_values = ols_model.pvalues
```

```
print(f"RSS: {RSS}")
```

```
print(f"RSE: {RSE}")
```

```
print(f"MSE: {MSE}")
```

```
print(f"R-squared: {R_squared}")
```

```
print(f"Standard Errors: {standard_errors}")
```

```
print(f"t-statistics: {t_statistics}")
```

```
print(f"p-values: {p_values}")
```

```
c:\Users\HP\AppData\Local\Programs\Python\Python311\Lib\site-packages\numpy\core\fromnumeric.py:86:
    return reduction(axis=axis, out=out, **passkwargs)
```

```
RSS: aveOralM      15.170504
```

```
dtype: float64
```

```
RSE: aveOralM      0.276104
```

```
dtype: float64
```

```
MSE: 0.07436521744807979
```

```
R-squared: 0.6468420800555861
```

```
Standard Errors: const      0.803926
```

```
T_OR1      0.883501
```

```

T_OR_Max1      0.882069
T_FHC_Max1     0.044464
T_FH_Max1      0.049258
dtype: float64
t-statistics:  const      8.753146
T_OR1          0.104128
T_OR_Max1      0.526115
T_FHC_Max1     -1.964102
T_FH_Max1      7.529419
dtype: float64
p-values:  const      1.191574e-17
T_OR1       9.170938e-01
T_OR_Max1   5.989521e-01
T_FHC_Max1  4.985945e-02
T_FH_Max1   1.358512e-13
dtype: float64

```

11. Significant and Insignificant features

In linear regression, we consider a feature significant if its p-value is less than 0.05. Conversely, if the p-value is greater than or equal to 0.05, we regard the feature as insignificant.

```

significant_features = p_values[p_values < 0.05]
insignificant_features = p_values[p_values >= 0.05]

print(f"Significant Features: {significant_features}")
print(f"Insignificant Features: {insignificant_features}")

Significant Features:  const      1.191574e-17
T_FHC_Max1      4.985945e-02
T_FH_Max1       1.358512e-13
dtype: float64
Insignificant Features: T_OR1      0.917094
T_OR_Max1       0.598952
dtype: float64

```


4 Performance evaluation of Linear regression

2. Residual Standard Error (RSE)

The Residual Standard Error (RSE):

$$\text{RSE} = \sqrt{\frac{\text{SSE}}{N - d - 1}}$$

N = Total number of data samples

d = The number of independent features

Model A:

$$\text{RSE}_A = \sqrt{\frac{9}{10000 - 2 - 1}} \approx \sqrt{\frac{9}{9997}} \approx 0.03$$

Model B:

$$\text{RSE}_B = \sqrt{\frac{2}{10000 - 4 - 1}} \approx \sqrt{\frac{2}{9995}} \approx 0.01$$

- Since Model B has a lower RSE, Model B fits more with the dataset.

3. R-squared (R^2)

$$R^2 = 1 - \frac{\text{SSE}}{\text{TSS}}$$

Model A:

$$R_A^2 = 1 - \frac{9}{90} = 1 - 0.1 = 0.9$$

Model B:

$$R_B^2 = 1 - \frac{2}{10} = 1 - 0.2 = 0.8$$

- Model A has a higher R^2 , indicating it explains more variance in the response variable.

4. Metrics Comparison

1. Scale Independence:

- **R^2 :** R^2 is a unitless measure that indicates the proportion of variance in the dependent variable that is explained by the model. This makes it scale-independent, meaning it remains consistent regardless of the range or units of the data. This property allows for fair comparisons between models across different datasets or variables with varying scales.
- **RSE:** RSE is measured in the same units as the dependent variable, so its value can vary depending on the scale of the data. This means RSE's value is influenced by the range of the dataset, making it less straightforward to compare models if the datasets have different units or scales.

2. Normalized Benchmark for Model Performance:

- **R^2 :** Since R^2 ranges from 0 to 1, it provides a normalized benchmark for evaluating how well a model explains the variance in the data. A higher R^2 value indicates better model performance, making it easier to assess and compare models directly.
- **RSE:** Although RSE indicates the average size of residuals, its absolute value can be influenced by the data's scale and units. This makes it harder to compare models across different datasets, as the RSE values are not normalized and can vary with the scale of the outcome variable.

5 Linear regression impact on outliers

2. What happens when $a \rightarrow 0$?

When a approaches 0, both modified loss functions $L_1(w)$ and $L_2(w)$ change their behavior significantly. Let's consider each function:

- **For $L_1(w)$:**

$$L_1(w) = \frac{1}{N} \sum_{i=1}^N \left(\frac{r_i^2}{a^2 + r_i^2} \right)$$

As a approaches 0, the term a^2 becomes negligible compared to r_i^2 , so:

$$L_1(w) \approx \frac{1}{N} \sum_{i=1}^N \left(\frac{r_i^2}{r_i^2} \right) = \frac{1}{N} \sum_{i=1}^N 1 = 1$$

This implies that $L_1(w)$ converges to 1 for every data point, making the loss function independent of the residuals r_i . Essentially, the influence of outliers becomes uniform across all data points.

- **For $L_2(w)$:**

$$L_2(w) = \frac{1}{N} \sum_{i=1}^N \left(1 - \exp \left(-\frac{2|r_i|}{a} \right) \right)$$

As a approaches 0, $\frac{2|r_i|}{a}$ becomes very large, so $\exp \left(-\frac{2|r_i|}{a} \right)$ approaches 0. Thus:

$$L_2(w) \approx \frac{1}{N} \sum_{i=1}^N (1 - 0) = 1$$

Similar to $L_1(w)$, $L_2(w)$ also converges to 1 for all data points, meaning that all residuals are treated the same regardless of their size.

3. Minimizing the influence of data points with $|r_i| \geq 40$

To minimize the influence of data points with $|r_i| \geq 40$, we need to choose a value of a and a loss function that allows us to easily identify points where $|r_i| \geq 40$.

Thus, we should select a loss function and an a value such that the loss function increases significantly for $|r_i| \geq 40$ and remains relatively smaller for $|r_i| < 40$.

In my opinion, the L_1 loss function with $a = 25$ is a better choice for this purpose.

```
import numpy as np
import matplotlib.pyplot as plt

# Define the range of r_i values
r_i = np.linspace(-50, 50, 500)

# Define the values of a
a_values = [2.5, 25, 100]

# Define the L1 and L2 loss functions
def L1(r_i, a):
    return r_i**2 / (a**2 + r_i**2)

def L2(r_i, a):
    return 1 - np.exp(-2 * np.abs(r_i) / a)

# Plot L1 and L2 Loss Functions
plt.figure(figsize=(14, 7))
```

```

# Plot L1 Loss Function for different values of a
for a in a_values:
    plt.plot(r_i, L1(r_i, a), label=f'L1 Loss with a={a}')

# Plot L2 Loss Function for different values of a
for a in a_values:
    plt.plot(r_i, L2(r_i, a), linestyle='--', label=f'L2 Loss with a={a}')

# Add labels and legend
plt.xlabel('$r_i$')
plt.ylabel('Loss')
plt.title('Comparison of Loss Functions: $L_1$ and $L_2$')
plt.legend()
plt.grid(True)

# Show the plot
plt.show()

```

